

Correlation among smart grid actors

Nicolas Gensollen, Monique Becker, Vincent Gauthier and Michel Marot

CNRS UMR 5157 SAMOVAR,

Telecom SudParis/Institut Mines Telecom

Email: {nicolas.gensollen, vincent.gauthier, michel.marot, monique.becker}@telecom-sudparis.eu

Abstract—Achieving a successful energetic transition became a major goal for the 21st century in most developed countries. The reconciliation of several research areas and their recent focus on the same smart grid objective has led to major technical improvements. For instance, allowing a large penetration of renewables into the electricity production grid seemed more like a dream yesterday, but appears more and more doable as research progresses. The inherent volatility of some of these energy sources poses indeed various problems in such a constrained environment. It is often supposed that a large number of these small generators will be decentralized down to the distribution network and even be owned by individuals. By enabling bi-directional flows, this vision revolutionises the nature of today simple consumers into futuristic agents that can both consume and produce. Such agents can obviously reduce their electricity bills by seeking energetic independence through demand side management programs coupled with storage. Nevertheless, as it could be the case that an agent's production exceeds its consumption and storage capacity, its ability to sell its extra-production appears as a way to save energy.

In this paper, we thus consider agents that tend to produce more than they consume and seek to sell their surplus of production. More precisely, we rely on aggregations of these agents that can be seen as virtual power plants with production and consumption components that are both variable. Because these virtual power plants are complex aggregations of different types of production means and consumption habits, the volatilities of the plants productions are deeply impacted by the correlation between the agents. We propose here an algorithm that takes the underlying correlation structure of the agents into account while forming the coalitions. We show then that the resulting diversified coalitions are able to generate higher benefits on a constrained energy market, and are more resilient to random failures of the agents.

I. INTRODUCTION

Designing stable power systems is a classical engineering challenge since blackouts can have catastrophic consequences. One obvious condition for stability revolves around sustaining at any time a power generation that meets the demand. If one is larger than the other, the system deviates from its synchronous state. If nothing is done to return to the synchronized equilibrium, this can lead to catastrophic cascades of failures. With traditional power plants based on fossil energies, the production can be scheduled in order to sustain the predicted consumption. Deviations from this schedule can then be supported almost on real time by using fast response power plants, interconnections with border countries, as well as regulator entities. Even if most individual consumers may not realize it, the electricity prices are not constant rates and evolve with the production/consumption conditions on the

grid. It is often assumed that the granularity of these prices is meant to increase in the future as a way to reverberate the production conditions on the end users. The pricing of electricity and the necessity for the grid operator to have different kinds of emergency reserves lead to an economy setting for electricity, where operations and various kind of contracts are decided on a market. Such a market environment clearly necessitates communication in order to monitor and obtain information from the grid as well as exchange between participants. These kind of settings are already being used at the transport level with large renewable plants in the production portfolio, and they appear as a way to manage the reliability of the whole system, especially when the penetration of renewables is important.

The smart grid vision however goes further and revolutionizes this top-down centralized architecture by assuming that bidirectional power flows are allowed. This would change completely the nature of the traditional grid since production could be located down to the very end of the distribution networks. Nodes that were simply pure loads yesterday could behave tomorrow as generators or loads depending on time. On the other hand, the penetration of renewables in the production is continuously growing, and is expected to become majority in a near future. These plants completely rely on the presence and the intensity of their respective resources (wind, sun, tides...), which predictions can be difficult and uncertain. Balancing production and demand in such a scenario appears much more challenging and it seems inevitable that the grid should modernize its infrastructures to sustain this transition. The key to drive safely such a system is assumed to be information, and more precisely, the capacity to measure, communicate, and analyze data on real time.

In this paper, we focus on these "*nodes*" in the distribution network that can produce and consume electricity. More precisely, we consider a set of agents that can be seen as aggregators of renewable distributed energy resources (DER) and electrical loads. The production of an agent can be, of course, used to meet its own demand, but in cases where it is over-producing, we consider that it has the possibility to inject, against financial compensation, its extra-production in the grid. Such an agent model is known as a "*prosumer*" [7] (and will be called accordingly in this paper).

As beautiful as this distributed participatory production scenario seems to be, the constraints necessary for the electrical grid to remain in a stable state obviously forbid any anarchic system where prosumers inject power whenever and however

they want. Actually the architecture and control of such a distributed grid creates numerous interdisciplinary challenges. In this paper, the envisaged solution for organizing such a multi-producers system consists in extending the kind of market environment discussed in the previous section. Participating aggregators estimate a production capacity for an upcoming period of time and announce it on the market. If a contract between the grid and a participant is bought, the latter commits to injecting exactly, at any time of the contracted period, the announced amount of power, under financial penalties if it fails.

In this paper, we consider that the production of a prosumer comes from renewables. Contracts are therefore based on predictions and come with a given amount of risk due to the intermittent nature of most renewables. Using storage as buffers in order to reduce risks is a popular approach. Actually, a whole branch of the smart grid literature is even studying the possibility to use electric vehicles as moving storage capacities for stabilizing the grid. Without storage and proper control, the over-producing state of the agents, and thus the power they inject in the grid, might be rather unstable. This is clearly unacceptable for the grid operator that cannot ensure system stability if it has to deal with numerous small unpredictable entities.

Forming coalitions/aggregations is an envisaged solution to both the number of entities soliciting the operator, and their high variabilities. Indeed, it has been argued that using a multi-level aggregators architecture for the control could lead to better performances on the communication side, since an aggregator can be considered as a single node by the level above it. On the other hand, it is well-known that diversification of the assets is a way of minimizing risks when constructing a portfolio. One thus expects a more stable and predictable energy production for a coalition than for a single agent. Nevertheless, all coalitions are obviously not equally stable, which means that special attention should be paid to the aggregation step.

Furthermore, in our settings, since the prosumers have both a consumption and a production component, each depending on location and time, the diversification of the aggregations seems not trivial. This is a central topic of the present paper : given N prosumers, what coalitions should be formed so that the compromise between expected production and variability is optimized ?

We will see that variability in the aggregated productions can be quantified to a certain extent by the correlation among the agents forming the coalitions. Understanding the correlation relationships among the agents can thus be illuminating in the sense that it gives an indication about what coalitions to form and how much they should sell. More precisely, we use a graph representation of the correlation structure to gain insight about the expected production to risk ratios of different coalitions. We build a framework in which the system operator has the possibility of specifying both a maximum amount of risk and a minimum production acceptable to enter the market. Based on predictions, coalitions can decide whether or not to

enter and what power quantities they are willing to provide. Of course, coalitions failing at fulfilling their obligations during the contracted period of time (because of bad forecasts, unexpected rare events, or intentional false information) will be exposed to financial penalties.

Because agents are susceptible to fail for diverse reasons, the impact on the system's stability is critical. The propensity of a system to undertake these failures is often referred to as its resilience [3]. Once formed, we also investigate in this paper the resilience of the coalitions when prosumers fail. Despite the fact that losing agents is usually detrimental to the coalitions, we will see that, with our algorithm, coalitions tend to be less impacted by random failures.

The paper is organized as follows, section II gives a brief overview of the related literature, section III clarifies how we generated realistic prosumer production traces based on weather data. In section IV, we define most of the notations and explain why correlation between prosumers is a quantity of interest for our objective. Based on the conclusions of section IV, we develop in section V a utility function quantifying how much power a coalition can announce on the market given an accepted risk level. We then develop, in section VI a greedy optimization algorithm that uses decorrelated cliques as inputs and improves their utility over a correlation-constrained environment. Finally, section VII provides some results both on performance of the method and resilience of the coalitions formed.

II. RELATED WORK

The intermittent nature of renewables such as wind or solar power introduces new challenges in control design. On the contrary to fossil plants whose production can be scheduled in advance to meet the expected consumption, renewables by definition only produce when the natural resource is present. Unfortunately, these moments do not necessarily coincide with the consumption peak hours [8]. A first approach to remedy this situation by acting on the consumption component is the use of dynamic electricity prices sent to the end users. Together with demand side management techniques implemented on the smart meters, they constitute a first tool for shaping the consumption profiles [8]. In periods when the production is expected to be greater than the consumption, prices are scheduled low in order to encourage end users to delay their elastic loads to these periods. On the contrary, when the production is expected to be less than the consumption, high prices are scheduled to deter any delayable loads over these periods.

Another popular approach on the production side, is to combine renewable generators with storage devices, such that these are charged when there is a surplus of production, and discharged when the consumption exceeds the production. Although simple, this idea causes numerous challenges. Depending on the size of the system considered, centralized or decentralized control is desirable. In [], the authors introduce a distributed energy management system with a high penetration of renewables such that power is scheduled in a distributed

fashion. In [] the optimal storage capacity problem is addressed. There is indeed an interesting tradeoff between the costs of the equipments and the expected availability of power. The authors develop a framework that enables them to exhibit a Pareto front of efficient solutions.

Although there are great progresses in stabilizing the production by demand side management or storage architectures, the need for good prediction techniques is more important than ever. Understanding the statistical properties of wind or solar irradiance are already well-established areas of research. There exists indeed a very interesting literature on techniques aiming at predicting wind or solar irradiance profiles for future periods with error margins. In this paper, we adopt a slightly different approach : we suppose that the agents have historical records of their production and consumption, and we seek aggregations of agents with high expected production and low variability.

Optimization of expected returns to risk is a traditional goal in finance. It is indeed well-known, that the more risk one is willing to take, the higher his potential gains. On the contrary, when investing exclusively on low risk assets, one should expect relatively small gains. This tradeoff is formalized in the Markowitz' portfolio theory [9]. More precisely, given a set of assets for which we have some historic data of returns, the objective is to find a linear combination of these assets (the so-called portfolio) which maximizes the expected value while minimizing the variance of the portfolio's return. Markowitz's answer is a set of efficient portfolios that all optimize in some sense this tradeoff. If one is able to put a number on his risk acceptance or on the target expected return, the corresponding efficient portfolio is a priori the best option.

This problem exhibits similarities with our objective of forming stable coalitions since our expected return can be formulated as the expected production and the risk / volatility as the variance of the production. However, it is assumed in the portfolio theory [9] that returns are normally distributed. In this paper, we chose to use normal distributions only for explanation purposes, or for estimating some parameters. Therefore, all the results are computed based on observed distributions (sometimes far from gaussians).

One of the key point in the Markowitz theory is to consider explicetely the correlation structure between the assets since these correlation relationships impact directly the variances of the portfolios. Since the work of [10], an interesting approach consists in computing a distance metric based on the correlation coefficients in order to organize the series in a correlation graph where the weight of an edge between two series (nodes) is the metric value for these series.

Because the metric can be computed for all pairs, these graphs are complete and of little use as is. Historically, the approach used by [10] was to compute a minimum spanning tree over the correlation graph as to extract a correlation structure of the form of a hierarchical clustering. Later on, it was pointed out that, by definition, a spanning tree could not capture the underlying clustering structure hidden in the correlation graph. Another filtering approach by mean of a threshold ϵ on the edge weights solves this problem and allows

one to exhibit clusters of correlated series. In the following we will refer to these filtered graphs as ϵ -graph.

An efficient portfolio of assets and a clustering of assets based on the correlations are obviously two very distinct things. And none of them is exactly what we seek. We will see in section X, how we combined both theories in order to form sufficiently producing coalitions with an acceptable volatility level.

III. GENERATING REALISTIC PROSUMER PATTERNS

An essential component of the smart grid is the smart meter which makes the interface between the end user and the rest of the system. Smart meters coupled with sensors measure quantities of interest (like instantaneous consumption), receive informations from the grid (electricity price for instance), and take actions accordingly (demand side management program). Smart meters are currently and gradually deployed, and will probably provide interesting datasets to work on. Unfortunately, at the time this paper was written, production and consumption data for prosumers over a large region were not yet available to our knowledge. Some interesting experiments are notwithstanding being conducted and data are progressively made public [].

In the following, we denote by $P_i(t)$ the instantaneous extra-production of agent i at time t :

$$P_i(t) = P_i^P(t) - P_i^D(t) \quad (1)$$

Where $P_i^P(t)$ represents the total production of agent i at time t and $P_i^D(t)$ its consumption at time t . In other words, $P_i(t)$ represents the instantaneous surplus of power that agent i is willing to sell at time t . As explained above, since large datasets containing this quantity over time are not yet available, we simulated these traces by considering separately P_i^P and P_i^D .

For a prosumer i , it is possible to write both quantities as a sum over the distributed energy resources (DER_i) and loads ($load_i$) of i :

$$P_i^P(t) = \sum_{k \in DER_i} P_k(t) \quad (2)$$

$$P_i^D(t) = \sum_{k \in load_i} P_k(t) \quad (3)$$

For simplicity, in this paper we only consider windturbines (WT) and photovoltaic panels (PV) as possible DER for the agents ($DER_i = WT_i \cup PV_i$):

$$P_i^P(t) = \sum_{k \in WT_i} P_k(t) + \sum_{k \in PV_i} P_k(t) \quad (4)$$

We denote by $\nu_i(t)$ and $\xi_i(t)$ the wind speed (in $m.s^{-1}$) and the solar irradiance (in $W.m^{-2}$) at the location of agent i and at time t , so that :

$$P_i^P(t) = \sum_{k \in WT_i} \mathcal{F}_{WT}(\nu_i(t)) + \sum_{k \in PV_i} \mathcal{F}_{PV}(\xi_i(t)) \quad (5)$$

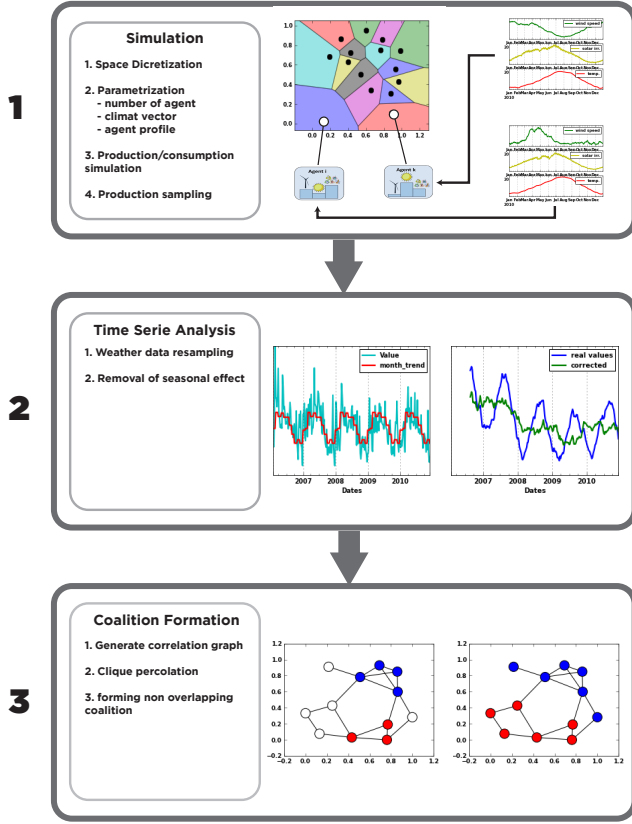


Figure 1: Process diagram

Where \mathcal{F}_{WT} (resp. \mathcal{F}_{PV}) is the power curve for the wind-turbines (resp. photovoltaic panels). We made here the implicit assumption that all windturbines (resp. photovoltaic panels) have the same power curve. The model can be easily extended to multiple power curves accounting for different types of generators. More details about power curves and their approximations can be found in [].

Weather quantities like wind speed or solar irradiance appear thus as alternative data for generating the P_i series. Fortunately, these kinds of data are easier to find, and since the development of small personal weather stations, their geographical granularity keeps increasing. A quantity like wind speed depends both on time and location (we neglect height since this information is rarely present in the data). In the following, time is discretized into slots and space into zones. A zone is simply a portion of the considered space for which we have weather data. Therefore, if prosumers i and j are positioned on the same zone, they are exposed to the same weather. Adding some intra-zone noise can easily be done though not considered in this paper. The process for generating the P_i series is pictured in the first block of the process diagram (see figure 1).

Note that a prosumer i is defined by his zone Z_i as well as the sets DER_i and $load_i$. That is, a prosumer can be configured to represent anything from a single windturbine for instance ($DER_i = \{WT_0\}$ and $load_i = \emptyset$) to a pure load ($DER_i = \emptyset$ and $load_i = \{L_0\}$) through more complex combinations. In practice, we use random configurations for the agents.

In the rest of the paper, we use french weather data (see www.infoclimat.fr) starting in january 2006 and ending in december 2012, with a sampling frequency of three hours, and generate N timeseries of extra-production over this date range.

IV. NOTATIONS

This section provides most of the notations and introduces important concepts for the rest of the paper. As explained in section III, we consider a set $\mathcal{A} = \{a_1, a_2, \dots, a_N\}$ of N prosumers configured randomly, and for each agent, we simulate its extra-production $P_i(t)$, $\forall i \in \mathcal{A}$ from 2006 to 2012. Based on these historical values, our objective is now to form groups of prosumers (the so-called coalitions) so that the global power production resulting from the superposition of individual's extra-productions be both sufficiently high and predictable. Let $P_S(t) = \sum_{i \in S} P_i(t)$ be the extra-production of coalition S at time t .

Suppose now that coalition S has to suggest a production value P_S^{CRCT} to enter the market. This means that, during the time S is on the market, it will have to inject in the grid exactly P_S^{CRCT} at any time t and will be rewarded proportionally to this amount, with penalties if it deviates. Obviously, the actual extra-production will not be constant at this value and will oscillate due to intermittencies in the production and consumption. If S always produces more than P_S^{CRCT} , it will never have to pay penalties, but it is losing some gains since it could have announced a higher contract value. If the production oscillates around P_S^{CRCT} , by using batteries or demand side management techniques (see section II), S could be able to maintain its production to the contract value at any time. Nevertheless, if the oscillations are too important compared to the available storage capacity, S will probably break the contract and pay penalties. We can see that there is a return over risk tradeoff here, meaning that coalitions should find the right balance between announcing too low and losing some potential gains, and claiming too high and paying penalties.

Let us illustrate the rest of the notations and concepts with a simple example. We consider only two agents i and j such that the distribution of their extra-production can be approximated by normal distributions: $P_i \sim \mathcal{N}(\mu_i, \sigma_i)$ and $P_j \sim \mathcal{N}(\mu_j, \sigma_j)$. This is only for explanation purposes as it is of course rather unrealistic in real situations where the distributions are skewed. Using simple statistics, we can write the distribution of the coalition $S = \{i, j\}$ as $P_{\{i,j\}} \sim \mathcal{N}(\mu_{ij}, \sigma_{ij})$, where :

$$\begin{cases} \mu_{ij} = \mu_i + \mu_j \\ \sigma_{ij} = \sqrt{\sigma_i^2 + \sigma_j^2 + \rho_{ij}\sigma_i\sigma_j} \end{cases} \quad (6)$$

ρ_{ij} being the Pearson's correlation coefficient between P_i and P_j . If the coalition $\{i, j\}$ proposes a contract value P_S^{CRCT} , all instants where $\{i, j\}$ will produce less than P_S^{CRCT} is critical. Indeed, in this kind of situations, $\{i, j\}$ will either have to discharge batteries to keep up with its contract, or pay penalties to the grid. The probability that $\{i, j\}$ is underproducing compared to the contract : $Pr[P_{i,j} \leq P_S^{CRCT}]$ is thus an important indicator of the coalition's quality. A well-known result for normal distributions is that the cumulative distribution function can be written as :

$$Pr[P_{ij} \leq P_S^{CRCT}] = \frac{1}{2} \left[1 + \operatorname{erf} \left(\frac{P_S^{CRCT} - \mu_{ij}}{\sigma_{ij} \sqrt{2}} \right) \right] \quad (7)$$

Where erf is the error function : $\operatorname{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$.

The amount of risks a given coalition is willing to take depends on a lot of things, among which its capacity to compensate for under-producing (using batteries, backup generators...), or the risk aversion of the aggregator. Selecting the right contract value appears thus as an interesting problem on its own that we plan to investigate in future works. In order to keep the present paper in a reasonable length, we simplify the contract value selection problem a little bit by giving some responsibilities to a third party named the grid operator. The role of the grid operator is to constrain the market entry to coalitions able to propose both sufficiently high and sufficiently credible contract values. More formally, let $\phi \in [0, 1]$ be the reliability threshold fixed by the grid operator as a maximum value for the probability of under-producing. The highest contract value that a coalition can propose is thus P_S^{CRCT*} such that $Pr[P_{ij} \leq P_S^{CRCT*}] = \phi$. In the gaussian example, this translates by coalition $\{i, j\}$ announcing :

$$P_S^{CRCT*} = \mu_{ij} - \sqrt{2} \sigma_{ij} \operatorname{erf}^{-1}(1 - 2\phi) \quad (8)$$

This is the best contract value that the coalition S can afford giving the stability policy ϕ of the grid operator. Figure 2 shows how P_S^{CRCT*} evolves according to the reliability parameter ϕ . For illustration, the range of ϕ values is shown from 0 to 1, but in practice, only small values of ϕ really make sense : $\phi = 1$ for instance means that coalitions can announce absolutely anything since the probability of producing less than any contract value is necessarily less than one by trivial definition of a probability. As visible on figure 2, coalitions with high expected productions but presenting a high unpredictability are penalized and can only afford small contracts.

In order not to overload the market with unrealistically small coalitions, the grid operator also specifies a lower bound P^{MIN} on the contract values. We thus characterized a valid coalition as one satisfying the two conditions :

$$\left\{ \begin{array}{l} Pr[P_{ij} \leq P_S^{CRCT}] \leq \phi \\ P_S^{CRCT} \geq P^{MIN} \end{array} \right. \quad (9)$$

On figure 2, P^{MIN} is fixed to 2 units for illustration purpose. For $\phi = 0.1$, only blue triangles and cyan circles coalitions are valid while red diamonds coalition is not.




Figure 2: P_S^{CRCT*} depending on reliability parameter ϕ for gaussian distributions (see equation 8). Blue curve with triangles stands for a coalition S with an expected production of 5 units and a standard deviation of 0.5. Under a grid policy of $\phi = 0.1$, it is able to announce a contract value of $P_S^{CRCT*} = 4.36$. The same coalition in term of expected production ($\mu = 5$), but with a higher variance ($\sigma = 1.5$, cyan curve with circles) can only afford a smaller contract value of $P_S^{CRCT*} = 3.07$. The red curve with diamonds stands for a coalition with a higher expected production ($\mu = 7$), but with a very high unpredictability ($\sigma = 5$). For low values of ϕ , this coalition is thus heavily penalized and can only afford a contract of 0.59 units. Under grid policy ($\phi = 0.1, P^{MIN} = 2$), this last coalition is thus not allowed to enter the market (red dot below the horizontal dashed line).

The gaussian assumption of this small example is convenient as it allows us to write P_S^{CRCT*} analytically. Nevertheless, such assumption is rather unrealistic in practice. In the following, we keep the same framework but release this gaussian assumption unless the contrary is specified. This assumption will indeed be convenient for computing some parameter estimates.

V. UTILITY FUNCTION

In this section, we use the notions of contract values and valid coalitions developed in section IV in order to design a proper utility function. The contract basically indicates the rate at which a coalition has to inject power in the grid. It seems then natural that coalitions are remunerated proportionally to their contract values $gain(S) \propto P_S^{CRCT*}$. More precisely, if λ is the unitary price rate for electricity, a coalition S injecting P_S^{CRCT*} in the grid during a period $[t_0, t_k]$ earns :

$$gain(S) = \int_{t_0}^{t_k} \lambda P_S^{CRCT*} dt = P_S^{CRCT*} \int_{t_0}^{t_k} \lambda dt \quad (10)$$

(since P_S^{CRCT*} is supposed to be a constant rate over the contracted period). Using $gain(S)$ directly as a utility function suffers a major drawback. It is indeed not a concave function of

the coalitions' sizes, meaning that coalitions can grow as large as the number of agents allows it, without any counterparty.

Such a model, that virtually allows infinitely large coalitions and contract values, is in practice not realistic. There are indeed costs (communication costs for instance) that increase with the coalitions sizes. We take this observation into account by rescaling the utility of a coalition S by its size in term of number of agents ($|S|$):

$$\mathcal{U}(S) = \begin{cases} \frac{1}{|S|^\alpha} \frac{P_S^{CRCT*}}{P^{MAX}}, & \text{if } S \text{ is valid,} \\ 0, & \text{if } S \text{ is not valid} \end{cases} \quad (11)$$

Where parameter α controls to what extent the size of a coalition impacts its utility, and P^{MAX} is a normalizing factor. P^{MAX} can be seen as the maximum production which can be injected in the grid.

Based on \mathcal{U} , the marginal contribution of an agent i can be expressed as $\delta_S(i) = \mathcal{U}(S + \{i\}) - \mathcal{U}(S)$. A coalition S has thus an interest in adding an additional agent i if this marginal contribution is positive :

$$\delta_S(i) \geq 0 \Leftrightarrow P_{S+\{i\}}^{CRCT*} \geq P_S^{CRCT*} \left(\frac{|S|+1}{|S|} \right)^\alpha \quad (12)$$

If α is set to zero, agents are added as long as they increase the contract value of the coalition. If α is greater than zero, additional agents have to increase the contract value by some factor. Since tuning the utility function with α is not necessarily convenient, we relate α to the mean sizes of the coalitions \bar{N} . If we wish that $\mathcal{U}(S)$ tends to form coalitions of size approximately of the order of \bar{N} , then :

$$\left[\frac{\partial \mathcal{U}}{\partial |S|} \right]_{|S|=\bar{N}} = 0 \quad (13)$$

In order to get an estimator for α , we solve equation 13 in a gaussian case (as in section IV). Furthermore, since considering all the possible interactions between agents is analytically intractable, we use here a mean approximation. Any quantity x that varies over the agent set is thus simplified in its mean value \bar{x} . Solving equation 13 for α in these conditions yields :

$$\alpha_{\bar{N}}^* = \frac{0.7\bar{\sigma}(\bar{\rho}-1)\text{erf}^{-1}(2\phi-1)}{\bar{\mu}\sqrt{\bar{N}(\bar{\rho}\bar{N}-\bar{\rho}+1)} + 1.4\bar{\sigma}\text{erf}^{-1}(2\phi-1)(\bar{\rho}\bar{N}-\bar{\rho}+1)} \quad (14)$$

Figure 3 shows how α^* and the utility function evolves according to the mean size of the coalitions \bar{N} . These nice curves are of course only valid in the simplified example considered here, but they will provide some guidance when using real data.

As can be pointed out, the purpose of \mathcal{U} is not a study of coalitions stability against player defection, which game theory provides a lot of tools for. This is indeed a problem on its own. The redistribution of a coalition's utility in terms of individual payoffs will therefore not be considered directly in this paper. \mathcal{U} can be rather interpreted as a measure of how good a given




Figure 3: Gaussian mean approximation. Subplot **a** shows how the parameter α of the utility function should be chosen in function of the mean desired size of the coalitions (see equation 14). Subplot **b** displays the corresponding utility functions for different values of α . Blue curve with diamonds favors very small coalitions of 2 agents while the green one with triangles favors 5 agents coalitions. Finally, the red curve with squares has an optimal size of 15 agents.

coalition is according to our criteria since it favors coalitions with good production to risk ratios.

VI. COALITION FORMATION

Section IV explained how contract values for the coalitions are computed, and in section V we related this quantity to the utility and gains of a coalition. As the computation time of this quantity is not negligible, and is increasing in the sizes of the aggregations, we explain in this section the heuristic we used to compute the coalitions.

A. Representing the correlation structure

As seen in section IV, the variance of the aggregated production impacts directly the contract values, and depends on the covariances between the agents productions. We argue here that, by having some representation of the correlation structure between the agents, the search landscape for high utility coalitions could be reduced, such that good coalitions are more likely to be found quickly. Usually, this correlation structure is formalized with a covariance matrix or a correlation matrix that contains all the correlation coefficients between the agents : $M = (\rho_{ij})_{\forall i,j \in \mathcal{A}^2}$. By using a metric to map this matrix in a weighted adjacency matrix (see section II), it is possible to obtain a graph representation of the correlation relationships between the agents.

In the following, we use two opposite distance metrics for this mapping :

$$\begin{cases} d_{ij}^1 = 1 - \rho_{ij}^2, \\ d_{ij}^2 = \rho_{ij}^2 = 1 - d_{ij}^1 \end{cases} \quad (15)$$

Clearly, d^1 (resp. d^2) maps two correlated series as close points (resp. distant) while two uncorrelated series are distant (resp. close). These metrics enable us to compute a correlation graph $G_1 = (\mathcal{A}, E_1)$ and a "de-correlation" graph $G_2 = (\mathcal{A}, E_2)$. For any i and j , the weight of the edge e_{ij} is d_{ij}^1 in G_1 and d_{ij}^2 in G_2 .

In both cases, we want to keep only the edges which weights are located in the lower tail of the distance distributions. In other words, we want to compute the ϵ -graphs of G_1 and G_2 such that only meaningful edges remain. Selecting the right filter ϵ is thus an important point since it affects the landscape search for the coalition formation. Unfortunately, there seems to be no clear consensus in the litterature on how to select such a threshold. We will see later in this section that cliques in G_2 are potential seeds for the coalitions. Since we want to generate N_{COAL} coalitions, we need at least N_{COAL} cliques of a given size to start. Besides, since we consider coalitions as disjoint, the starting cliques should be non overlapping. We thus select our optimal threshold for G_2 as :

$$\epsilon^* = \min_{\epsilon \in [0,1]} \{ \epsilon \text{ s.t. } |\Theta_k(G_2^\epsilon)| \geq N_{COAL} \} \quad (16)$$

Where G_2^ϵ is the de-correlation graph G_2 filtered by ϵ , and $\Theta_k(G)$ is the set of non overlapping cliques of size k in a given graph G . In other words we select ϵ^* as the smallest threshold possible such that the filtered de-correlation graph contains at least N_{COAL} non overlapping cliques of size k . The existence of ϵ^* as defined in equation 16 is not guaranteed. The users has indeed to provide consistent values of N_{COAL} or k compared to the size of the agent population \mathcal{A} .

B. Cliques as de-correlated starting seeds

Correlation can be seen as cosine of angles in L^2 , hence even if there is no strict transivity relation for correlation, there is, to a certain extent, some partial notion of it. More precisely, if a , b , and c are three items such that $\rho_{ab} > \delta$ and $\rho_{bc} > \delta$, then we know, by the cosine addition formula¹, that $\rho_{ac} > 2\delta^2 - 1$. That is, if a and b are strongly correlated ($\rho_{ab} > 0.9$) and b and c are also strongly correlated ($\rho_{bc} > 0.9$), then there is a high probability for a and c of being strongly correlated ($\rho_{ac} > 0.62$).

Nevertheless de-correlation seems like a more complex concept than correlation in the sense that there is not even a partial notion of transivity when it comes to it. Therefore, the clustering coefficients of G_1^ϵ is much higher than the one of G_2^ϵ . This can be seen as another formulation of [11] on the structural roles of weak and strong links on financial correlation graphs. Strong links, accounting for strong correlation relationships, are responsible for the clustering, while weak links provide the connectivity between clusters. Searching for clusters in G_2^ϵ and hoping that this strategy will provide a nice coalition structure of internally uncorrelated coalitions seems thus pointless.

¹ $\cos(a+b) = \cos(a)\cos(b) - \sin(a)\sin(b)$




Figure 4: Histograms of utility values for coalitions of size 3. Red bars stands for cliques in the decorrelation graph, and blue bars represents all the other possible triplets. As visible, cliques tend to exhibit higher utilities than randomly selected coalitions.

Consider now a clique in G_2^ϵ , which is a complete subgraph of G_2^ϵ . This is indeed a structure of interest for our purpose. Since there is a link for every pairs of nodes, we know, by construction, that a clique has a mean correlation and a maximum correlation less than ϵ .

Figure 4 shows the distributions of the utility values for cliques of size 3 (triangles) in $G_2^{\epsilon^*}$ and for all the other possible triplets of agents. It is clearly visible that cliques tend to exhibit higher utilities because of their de-correlation property. Choosing cliques in $G_2^{\epsilon^*}$ as coalitions seems therefore appealing. Nevertheless, the quality of the results seems to decrease as the sizes of the cliques increase. Indeed, the larger the desired cliques, the more dense $G_2^{\epsilon^*}$ becomes (see equation 16). There is a point where cliques results more from noisy edges than true de-correlation relationships, which decreases the quality of the results.

Directly mapping cliques to coalitions by this de-correlation oriented approach is thus not sufficient. It is indeed possible that adding agents to these cliques has the combined effect of increasing the expected production while decreasing its stability. The question revolves around measuring the benefits of this production surplus compared to the disadvantage of a higher unstability. This can be quantified by the marginal benefit of equation 12.

C. Algorithm

The algorithm takes inputs from :

- **The agents** : historical series of available productions P_i ,
- **The grid operator** : market entrance policy (P^{MIN}, ϕ) ,
- **The "user"** : Number of desired coalitions N_{COAL} and size of starting cliques k .

The first steps consists in computing the de-correlation graph G_2 as well as the optimal threshold ϵ^* . Cliques of size k in $G_2^{\epsilon^*}$ are considered as coalition seeds. The

next step is a local greedy improvement over the landscape represented by $G_2^{\epsilon^*}$. Cliques add alternatively the node i^* in their neighborhood that yields the best marginal benefit $MAX_{i \in N(clique)} \delta_{clique}(i)$ where $N(clique)$ is the neighborhood of a given clique. This addition occurs only if i^* is not already involved in another coalition, and if $\delta_{clique}(i^*) \geq 0$, meaning that utilities are increasing. The algorithm stops when all nodes are distributed in a coalition or when the global utility stops increasing. See the details in algorithm 1.

Data: P_i series,
Grid policy (P^{MIN}, ϕ) ,
Desired number of coalitions N_{COAL} ,
size of starting cliques k
Result: $CS = \{S_1, \dots, S_{N_{COAL}}\}$
Compute $G_2^{\epsilon^*}$;
Find the N_{COAL} cliques in $G_2^{\epsilon^*}$;
while $\mathcal{U}(CS)$ is improving **do**
 for each clique do
 Find i^* ;
 if $\delta_{clique}(i^*) \geq 0$ **then**
 $clique \leftarrow clique \cup \{i^*\}$;
 end
 if $\exists j \in clique, s.t. \delta_{clique}(j) < 0$ **then**
 $clique \leftarrow clique - \{j\}$;
 end
 end
end
Algorithm 1: Local greedy optimization algorithm

Data: Agent set \mathcal{A} ,
Desired number of coalitions N_{COAL} ,
Maximum number of iterations $Loop^{max}$
Result: $CS = \{S_1, \dots, S_{N_{COAL}}\}$
 $N^{loop} \leftarrow 0$;
 $CS^* \leftarrow \emptyset$;
while $N^{loop} < Loop^{max}$ **do**
 $CS \leftarrow SelectRandomCS()$;
 if $\mathcal{U}(CS) > \mathcal{U}(CS^*)$ **then**
 $CS^* \leftarrow CS$;
 end
 $N^{loop} \leftarrow N^{loop} + 1$;
end
return CS^*

Algorithm 2: Random algorithm

VII. RESULTS

The algorithm presented in the previous section is supposed to generate a given number of coalitions that have good utilities. Nevertheless, as it comprises mainly of a greedy optimization based on local improvements, the probability that the algorithm finds the global optimal coalitions set is very low. Actually, it is not obvious that a strict optimum exists at all. Besides, there is, to our knowledge, no state of the art

Data: P_i series,
Desired number of coalitions N_{COAL} ,
search step size $\beta \ll 1$
Result: $CS = \{S_1, \dots, S_{N_{COAL}}\}$
 $\epsilon \leftarrow 1$;
 $CS \leftarrow \emptyset$;
while $|CS| < N_{COAL}$ **do**
 Compute G_1^ϵ ;
 $CS \leftarrow computeClusters(G_1^\epsilon)$;
 if $|CS| = N_{COAL}$ **then**
 return CS ;
 end
 else
 $\epsilon \leftarrow \epsilon - \beta$;
 end
end
Algorithm 3: Correlated algorithm

algorithm that aggregate uncorrelated agents in an optimum way (see section II for related problems).

In order to have an idea about the algorithm's quality, we compare its results with :

- **Random :** This algorithm is basically a random search over the coalition structures space. It analyses a given number of structures and returns the best that it has encountered so far. See algorithm 2.
- **Correlated :** This is the complete opposite of our algorithm. It uses the correlation graph G_1 and performs a community detection. The resulting coalitions have thus very high internal correlations. We thus expect this algorithm to perform very bad compared to the others. See algorithm 3.

Before running the algorithms, we need to calibrate the utility function by choosing the value of the α parameter. Recall that the purpose of this parameter is to take into account some constraints on the coalition's sizes if needed. In this paper, neither the communication network nor the electrical grid are explicitly considered. Thus, we do not have any technical constraints on coalition sizes even if we designed the utility such that these could be taken into account. We select the desired size as being $\lfloor N/N_{COAL} \rfloor$ (where $\lfloor \cdot \rfloor$ means floor). Figure 5 shows how the mean utility of a coalition evolves with its size when the optimum size is set to 40 agents. Using equation 14 to estimate α based on the mean quantities and gaussian approximations seems to give acceptable results for the utility function behavior.

Figure 6 displays the evolution of the global utility and the number of involved agents during the course of the greedy algorithm 1. The transition from invalid to valid coalitions is clearly visible on the blue diamond curve and occurs between iteration 10 and 15. After this transition, coalition's utilities improve slowly up to a maximum point.

Figure 7 shows the coalitions formed with the considered algorithms. A coalition is represented by a marker in the contract value / volatility space, and the color and shape of




Figure 5: Utility of random coalitions depending on their size. Blue dots shows real mean utility values and the thick red curve its smoothed version by applying a Savitzky-Golay filter. On this plot the α parameter of the utility function was selected according to equation 14 in order to favor 40 agents coalitions.




Figure 6: Evolution of the global utility $U(CS)$ (blue diamond curve, left axis) and the number of agents involved in the coalitions (red circle curve, right axis) during the greedy optimization of algorithm 1

a marker indicates by which algorithm the coalition has been formed. Besides, as the size of a marker is proportional to the number of agents in the coalition, we can see that the utility function results in approximatively balanced coalitions. The coalitions of correlated agents (green squares) are clearly of poor quality according to our criteria since they can only afford small production contracts, and with a very high volatility (compared to the other coalitions on the plot).

On figure 7, the de-correlated coalitions (blue dots) and the random ones (red triangles) are closer to the bottom right corner indicating a much better quality. We can see that, despite its simplicity, the random algorithm is able to achieve correct results thanks to the diversification inherent to such algorithm. Nevertheless, based on the utility, coalitions




Figure 7: Coalitions formed in the (contract value, volatility) space. The color map indicates qualities of portions of the plane. The closer to red the better (high contract values with small volatility). On the opposite, blue areas shows poor quality (small contract values with high volatility). Blue dots stand for the un-correlated coalitions, red triangles for random search, and green squares for correlated coalitions.

resulting from this random search are less good than the de-correlated ones. This can be understood on figure 7 as the de-correlated coalitions are globally more located toward the bottom right corner, meaning that there are less disparities between coalitions than with the random case.

A key point for the coalitions, besides stability and productivity, is their resilience. The resilience of a system can be roughly described as its ability to perform its tasks when subject to failures of its components. Therefore, the notion of resilience we will use in the following can be seen as the ability of the coalition structures to inject stable power in the grid when node failures occur. According to our model, the grid operator specified two thresholds (P^{MIN} and ϕ) such that the power injected by every coalition is constrained : $P_S^{CRCT} \in [P^{MIN}, P_S^{CRCT*}]$. As long as a coalition can propose a contract value higher than P^{MIN} , it is valid and allowed to enter the energy market. We define the resilience of a coalition S as the probability that S produces more than the P^{MIN} threshold :

$$\mathcal{R}_S = Pr[P_S \geq P^{MIN}] = 1 - Pr[P_S < P^{MIN}] \quad (17)$$

And we extend this measure to the coalition structures :

$$\mathcal{R}_{CS} = \prod_{S \in CS} (1 - Pr[P_S < P^{MIN}]) \quad (18)$$

We consider that prosumers fail randomly, and we denote by $\psi \in [0, 1]$ the fraction of agents that failed. Figure 8 exhibits how the resilience of the coalition structures evolve according to ψ . On the top subplot, P^{MIN} was voluntarily selected relatively low such that the resiliences of the three structures fit on the same figure. When the P^{MIN} requirement increases,

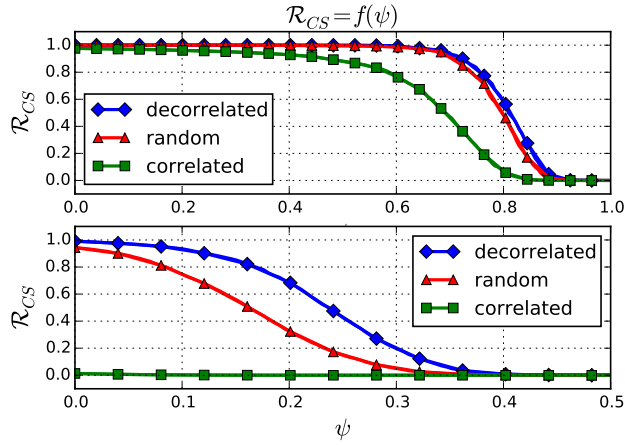


Figure 8: Resilience of the coalition structures when nodes fail randomly (see equation 18) for $P^{MIN} = 10MW$ (top subplot) and $P^{MIN} = 80MW$ (bottom subplot)

the differences between the algorithms also increase as visible on the bottom subplot of figure 8. The de-correlated coalitions seem to achieve a more resilient production on the market in the sense that they are able to sustain a higher fraction of node failures.

VIII. CONCLUSION

In this paper we studied how aggregations of prosumers could be authorized to sell their surplus of production to the grid operator. By relying on the past values of the agents, we constrained the market entry to both sufficiently productive and stable coalitions. The power that a coalition is able to propose on the market is therefore related to production and stability. As the correlations between the prosumers that form these coalitions impact directly their volatilities, we seek uncorrelated aggregations of agents. We used a graph representation of the correlation relationships between the agents as a reduced landscape for the coalition formation. A greedy algorithm that starts with cliques of the "de-correlation" graph of the agents and makes local improvements offers a good compromise between speed and quality of the results. We compare these results with a random search, and an opposite strategy that clusters correlated agents together. We showed that the coalitions resulting from our algorithm are able to provide more power to the grid with a lower volatility. Because they tend to have globally better production over volatility ratios, these coalitions will tend to use less storage and spill less energy than more unstable ones. We plan to study these benefits for the control of the aggregations in future works.

Because in real situations, agents are prone to failure, resilience is also an important criterium for the quality of the aggregations. We therefore studied how the coalitions are able to remain on the market when their agents fail randomly one by one. We showed that, in this situation, the coalitions resulting from our algorithm better withstand losses of agents.

REFERENCES

- [1] G. A. Pagani and M. Aiello, "The Power Grid as a Complex Network : a Survey," no. Section 2, pp. 1–34, 2012.
- [2] M. Rosas-Casals, S. Valverde, and R. V. Solé, "Topological Vulnerability of the European Power Grid Under Errors and Attacks," *International Journal of Bifurcation and Chaos*, vol. 17, no. 07, pp. 2465–2475, Jul. 2007. [Online]. Available: <http://www.worldscientific.com/doi/abs/10.1142/S0218127407018531>
- [3] S. Pahwa, C. Scoglio, and N. Schulz, "Topological Analysis and Mitigation Strategies for Cascading Failures in Power Grid Networks."
- [4] G. Filatrella, a. H. Nielsen, and N. F. Pedersen, "Analysis of a power grid using a Kuramoto-like model," *The European Physical Journal B*, vol. 61, no. 4, pp. 485–491, Mar. 2008. [Online]. Available: <http://www.springerlink.com/index/10.1140/epjb/e2008-00098-8>
- [5] F. Dörfler, M. Chertkov, and F. Bullo, "Synchronization in complex oscillator networks and smart grids." *Proceedings of the National Academy of Sciences of the United States of America*, vol. 110, no. 6, pp. 2005–10, Feb. 2013. [Online]. Available: <http://www.ncbi.nlm.nih.gov/pubmed/23319658>
- [6] S. D. Ramchurn *et al.*, "Putting the 'smarts' into the smart grid: A grand challenge for artificial intelligence," *Commun. ACM*, vol. 55, no. 4, pp. 86–97, Apr. 2012.
- [7] A. J. D. Rathnayaka *et al.*, "Prosumer management in socio-technical smart grid," in *Proceedings of the CUBE International Information Technology Conference*, 2012, pp. 483–489.
- [8] M. Milligan and B. Kirby, "Utilizing Load Response for Wind and Solar Integration and Power System Reliability," no. July, 2010.
- [9] H. M. Markowitz, "Portfolio Selection: Efficient Diversification of Investments," pp. 217–234, 1959. [Online]. Available: <http://www.informaworld.com/openurl?genre=article&doi=10.1080/00137917908965215&magic=crossref>
- [10] R. Mantegna, "Hierarchical structure in financial markets," *The European Physical Journal B*, vol. 11, no. 1, pp. 193–197, 1999.
- [11] A. Garas *et al.*, "The structural role of weak and strong links in a financial market network," *The European Physical Journal B*, vol. 63, no. 2, pp. 265–271, 2008.